

Assignment 2

August 20, 2020

You are currently looking at **version 1.2** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.

1 Assignment 2 - Pandas Introduction

All questions are weighted the same in this assignment. ## Part 1 The following code loads the olympics dataset (olympics.csv), which was derived from the Wikipedia entry on [All Time Olympic Games Medals](#), and does some basic data cleaning.

The columns are organized as # of Summer games, Summer medals, # of Winter games, Winter medals, total # number of games, total # of medals. Use this dataset to answer the questions below.

```
In [1]: import pandas as pd
```

```
df = pd.read_csv('olympics.csv', index_col=0, skiprows=1)
```

```
for col in df.columns:
    if col[:2]=='01':
        df.rename(columns={col:'Gold'+col[4:]}, inplace=True)
    if col[:2]=='02':
        df.rename(columns={col:'Silver'+col[4:]}, inplace=True)
    if col[:2]=='03':
        df.rename(columns={col:'Bronze'+col[4:]}, inplace=True)
    if col[:1]==' ':
        df.rename(columns={col:'#'+col[1:]}, inplace=True)
```

```
names_ids = df.index.str.split('\s\(') # split the index by '('
```

```
df.index = names_ids.str[0] # the [0] element is the country name (new index)
```

```
df['ID'] = names_ids.str[1].str[:3] # the [1] element is the abbreviation or ID (take first 3 characters)
```

```
df = df.drop('Totals')
```

```
df.head()
```

```
Out[1]:
```

	# Summer	Gold	Silver	Bronze	Total	# Winter	Gold.1	\
Afghanistan	13	0	0	2	2	0	0	
Algeria	12	5	2	8	15	3	0	
Argentina	23	18	24	28	70	18	0	
Armenia	5	1	2	9	12	6	0	
Australasia	2	3	4	5	12	0	0	

	Silver.1	Bronze.1	Total.1	# Games	Gold.2	Silver.2	Bronze.2	\
Afghanistan	0	0	0	13	0	0	2	
Algeria	0	0	0	15	5	2	8	
Argentina	0	0	0	41	18	24	28	
Armenia	0	0	0	11	1	2	9	
Australasia	0	0	0	2	3	4	5	

	Combined total	ID
Afghanistan	2	AFG
Algeria	15	ALG
Argentina	70	ARG
Armenia	12	ARM
Australasia	12	ANZ

1.0.1 Question 0 (Example)

What is the first country in df?

This function should return a Series.

```
In [2]: # You should write your whole answer within the function provided. The autograder will call
# this function and compare the return value against the correct solution value
def answer_zero():
    # This function returns the row for Afghanistan, which is a Series object. The assignment
    # question description will tell you the general format the autograder is expecting
    return df.iloc[0]

# You can examine what your function returns by calling it in the cell. If you have questions
# about the assignment formats, check out the discussion forums for any FAQs
answer_zero()
```

```
Out[2]: # Summer      13
        Gold          0
        Silver        0
        Bronze        2
        Total         2
        # Winter      0
        Gold.1        0
        Silver.1      0
        Bronze.1      0
        Total.1       0
        # Games      13
```

```

Gold.2          0
Silver.2        0
Bronze.2        2
Combined total   2
ID              AFG
Name: Afghanistan, dtype: object

```

1.0.2 Question 1

Which country has won the most gold medals in summer games?

This function should return a single string value.

```

In [3]: def answer_one():
        ObjA='Gold'
        return df[df[ObjA]==df[ObjA].max()].index[0]
        answer_one()

```

```
Out[3]: 'United States'
```

1.0.3 Question 2

Which country had the biggest difference between their summer and winter gold medal counts?

This function should return a single string value.

```

In [4]: def answer_two():
        df['ObjDif'] = abs(df['Gold'] - df['Gold.1'])
        ObjBig=df[df['ObjDif'] == df['ObjDif'].max()].index[0]
        return ObjBig
        answer_two()

```

```
Out[4]: 'United States'
```

1.0.4 Question 3

Which country has the biggest difference between their summer gold medal counts and winter gold medal counts relative to their total gold medal count?

$$\frac{\text{Summer Gold} - \text{Winter Gold}}{\text{Total Gold}}$$

Only include countries that have won at least 1 gold in both summer and winter.

This function should return a single string value.

```

In [5]: def answer_three():
        ObjData = df[(df['Gold'] > 0) & (df['Gold.1'] > 0)].copy()
        ObjData['ObjMedals']=ObjData['ObjDif']/ObjData['Gold.2']
        ObjAns=ObjData[ObjData['ObjMedals'] == max(ObjData['ObjMedals'])].index[0]
        return ObjAns
        answer_three()

```

```
Out[5]: 'Bulgaria'
```

1.0.5 Question 4

Write a function that creates a Series called "Points" which is a weighted value where each gold medal (Gold.2) counts for 3 points, silver medals (Silver.2) for 2 points, and bronze medals (Bronze.2) for 1 point. The function should return only the column (a Series object) which you created, with the country names as indices.

This function should return a Series named Points of length 146

```
In [6]: def answer_four():
        ObjGold=df['Gold.2']*3
        ObjSilver=df['Silver.2']*2
        ObjBronze=df['Bronze.2']*1
        df['Points']=ObjGold+ObjSilver+ObjBronze
        return df['Points']
        answer_four()
```

```
Out[6]: Afghanistan      2
        Algeria           27
        Argentina        130
        Armenia           16
        Australasia       22
        Australia        923
        Austria           569
        Azerbaijan        43
        Bahamas           24
        Bahrain            1
        Barbados           1
        Belarus           154
        Belgium           276
        Bermuda            1
        Bohemia            5
        Botswana           2
        Brazil            184
        British West Indies 2
        Bulgaria          411
        Burundi            3
        Cameroon           12
        Canada            846
        Chile              24
        China             1120
        Colombia           29
        Costa Rica         7
        Ivory Coast        2
        Croatia            67
        Cuba              420
        Cyprus             2
        ...
        Spain             268
        Sri Lanka          4
```

Sudan	2
Suriname	4
Sweden	1217
Switzerland	630
Syria	6
Chinese Taipei	32
Tajikistan	4
Tanzania	4
Thailand	44
Togo	1
Tonga	2
Trinidad and Tobago	27
Tunisia	19
Turkey	191
Uganda	14
Ukraine	220
United Arab Emirates	3
United States	5684
Uruguay	16
Uzbekistan	38
Venezuela	18
Vietnam	4
Virgin Islands	2
Yugoslavia	171
Independent Olympic Participants	4
Zambia	3
Zimbabwe	18
Mixed team	38

Name: Points, dtype: int64

1.1 Part 2

For the next set of questions, we will be using census data from the [United States Census Bureau](#). Counties are political and geographic subdivisions of states in the United States. This dataset contains population data for counties and states in the US from 2010 to 2015. [See this document](#) for a description of the variable names.

The census dataset (census.csv) should be loaded as census_df. Answer questions using this as appropriate.

1.1.1 Question 5

Which state has the most counties in it? (hint: consider the sumlevel key carefully! You'll need this for future questions too...)

This function should return a single string value.

```
In [7]: census_df = pd.read_csv('census.csv')
        census_df.head()
```

```

Out[7]:
SUMLEV REGION DIVISION STATE COUNTY STNAME CTYNAME \
0 40 3 6 1 0 Alabama Alabama
1 50 3 6 1 1 Alabama Autauga County
2 50 3 6 1 3 Alabama Baldwin County
3 50 3 6 1 5 Alabama Barbour County
4 50 3 6 1 7 Alabama Bibb County

CENSUS2010POP ESTIMATESBASE2010 POPESTIMATE2010 ... \
0 4779736 4780127 4785161 ...
1 54571 54571 54660 ...
2 182265 182265 183193 ...
3 27457 27457 27341 ...
4 22915 22919 22861 ...

RDOMESTICMIG2011 RDOMESTICMIG2012 RDOMESTICMIG2013 RDOMESTICMIG2014 \
0 0.002295 -0.193196 0.381066 0.582002
1 7.242091 -2.915927 -3.012349 2.265971
2 14.832960 17.647293 21.845705 19.243287
3 -4.728132 -2.500690 -7.056824 -3.904217
4 -5.527043 -5.068871 -6.201001 -0.177537

RDOMESTICMIG2015 RNETMIG2011 RNETMIG2012 RNETMIG2013 RNETMIG2014 \
0 -0.467369 1.030015 0.826644 1.383282 1.724718
1 -2.530799 7.606016 -2.626146 -2.722002 2.592270
2 17.197872 15.844176 18.559627 22.727626 20.317142
3 -10.543299 -4.874741 -2.758113 -7.167664 -3.978583
4 0.177258 -5.088389 -4.363636 -5.403729 0.754533

RNETMIG2015
0 0.712594
1 -2.187333
2 18.293499
3 -10.543299
4 1.107861

```

[5 rows x 100 columns]

```

In [8]: def answer_five():
        ObjStname=census_df.groupby('STNAME')['COUNTY'].count().idxmax()
        return ObjStname
        answer_five()

```

Out[8]: 'Texas'

1.1.2 Question 6

Only looking at the three most populous counties for each state, what are the three most populous states (in order of highest population to lowest population)? Use CENSUS2010POP.

This function should return a list of string values.

```
In [9]: def answer_six():
        ObjData=pd.DataFrame(census_df.where(
            census_df['SUMLEV']==50).groupby(['STNAME'])
            ['CENSUS2010POP'].nlargest(3))
        ObjData=ObjData.reset_index()
        ObjList=list(ObjData.groupby(['STNAME']).sum()
            ['CENSUS2010POP'].nlargest(3).index)
        return ObjList
        answer_six()
```

```
Out[9]: ['California', 'Texas', 'Illinois']
```

1.1.3 Question 7

Which county has had the largest absolute change in population within the period 2010-2015? (Hint: population values are stored in columns POPESTIMATE2010 through POPESTIMATE2015, you need to consider all six columns.)

e.g. If County Population in the 5 year period is 100, 120, 80, 105, 100, 130, then its largest change in the period would be $|130-80| = 50$.

This function should return a single string value.

```
In [10]: def answer_seven():
        ObjData=census_df[census_df.SUMLEV==50].copy()
        ObjData['ObjMax']=ObjData[['POPESTIMATE2010', 'POPESTIMATE2011',
            'POPESTIMATE2012', 'POPESTIMATE2013',
            'POPESTIMATE2014', 'POPESTIMATE2015']
            ].max(axis=1)
        ObjData['ObjMin']=ObjData[['POPESTIMATE2010', 'POPESTIMATE2011',
            'POPESTIMATE2012', 'POPESTIMATE2013',
            'POPESTIMATE2014', 'POPESTIMATE2015']
            ].min(axis=1)
        ObjData['ObjDif']=ObjData['ObjMax']-ObjData['ObjMin']
        ObjCountry=ObjData[ObjData['ObjDif']==
            ObjData['ObjDif'].max()].iloc[0]['CTYNAME']
        return ObjCountry
        answer_seven()
```

```
Out[10]: 'Harris County'
```

1.1.4 Question 8

In this datafile, the United States is broken up into four regions using the "REGION" column.

Create a query that finds the counties that belong to regions 1 or 2, whose name starts with 'Washington', and whose POPESTIMATE2015 was greater than their POPESTIMATE 2014.

This function should return a 5x2 DataFrame with the columns = ['STNAME', 'CTYNAME'] and the same index ID as the census_df (sorted ascending by index).

```
In [11]: def answer_eight():
        ObjAns=census_df[((census_df['REGION'] == 1)|
```

```

(census_df['REGION'] == 2))&
(census_df['CTYNAME'].str.startswith("Washington")
 & (census_df['POPESTIMATE2015'] >
    census_df['POPESTIMATE2014']
    ))].loc[:, ['STNAME', 'CTYNAME']]

return ObjAns
answer_eight()

```

```

Out[11]:
          STNAME          CTYNAME
896         Iowa  Washington County
1419    Minnesota  Washington County
2345  Pennsylvania  Washington County
2355   Rhode Island  Washington County
3163    Wisconsin  Washington County

```